

Tammy Jordan

PROFESSIONAL PROFILE

Senior Data/AI/ML professional, 19 years in. Track record of taking ambiguous problems through design, delivery and the unglamorous operational work that follows. I mentor deliberately and write things down. Core stack: Kafka, LangChain, Python.

EMPLOYMENT HISTORY

Senior Applied Scientist, Northbridge IT

2016–Present | Anthonyshire

- Reduced label noise by redesigning the annotation guidelines, improving validation accuracy by 70%.
- Trained and deployed LangChain models serving 5M predictions per day at under 110ms p95.
- Cut inference cost by 35% through quantisation, batching and autoscaling on Nginx.
- Established offline/online evaluation parity, catching 30 silent regressions before release.

Nova Tech Group — Data Scientist

2011–2016 | South Eric

- Migrated ad-hoc notebooks into versioned Spark pipelines with reproducible runs and tracked experiments.
- Developed an NLP classification system in TensorFlow reaching 82% F1, replacing a manual review queue.
- Mentored 8 analysts moving into modelling work, with weekly paper reading and code review.

Applied Scientist

Airbnb | 2009–2011 | Zacharyborough (Hybrid - 2 days on-site)

- Built an end-to-end training pipeline on Spark, cutting retraining from days to 10 hours.
- Wrote the team's model documentation and review standard, now required before any production launch.

Applied Scientist, Ridgeline Analytics

2007–2009 | Vargasstad (Hybrid - 2 days on-site)

- Ran A/B tests across 15 model variants; the winning variant added an estimated \$150k in annual revenue.
- Productionised demand forecasting models that reduced stock-outs by 45% across 2 regions.

Oracle — MLOps Engineer

2006–2007 | Schneiderchester (Remote (UK))

- Translated ambiguous business questions into a modelling roadmap with product and finance stakeholders.
- Lifted recommendation click-through by 20% through feature engineering and retraining cadence in Python.

EDUCATION AND TRAINING

Bachelor, Economics

Rasmussenmouth State University | 2006

TECHNICAL PROFICIENCIES

Kafka (17 yrs), LangChain (7 yrs), Python (15 yrs), TensorFlow (10 yrs), Docker (13 yrs), BigQuery (13 yrs), NumPy (6 yrs), Airflow (9 yrs), Snowflake (11 yrs), Scikit-learn (9 yrs), Spark (16 yrs), Jenkins (18 yrs), Confluence (14 yrs), Nginx (9 yrs), Next.js (7 yrs), PyTorch (4 yrs)

TRAINING & CERTIFICATIONS

Databricks Certified Data Engineer Associate — Databricks, 2025

Tableau Desktop Specialist — Tableau, 2024

References available upon request.