

Joshua Bean

Profile

17 years in Data/AI/ML, most of it spent on systems other teams depend on. Comfortable leading technical decisions, pushing back on scope when it matters, and carrying a pager for what I build. Primary tools: Kafka and NumPy.

Professional Experience

Clearwater Systems — Machine Learning Engineer

Dec 2012 – Jun 2026 | East Nicole (Remote (UK))

- Built an end-to-end training pipeline on PyTorch, cutting retraining from days to 40 hours.
- Lifted recommendation click-through by 18% through feature engineering and retraining cadence in NumPy.
- Established offline/online evaluation parity, catching 10 silent regressions before release.

Data Scientist

Ridgeline Analytics | May 2010 – Dec 2012 | Bakerside

- Wrote the team's model documentation and review standard, now required before any production launch.
- Translated ambiguous business questions into a modelling roadmap with product and finance stakeholders.

Shopify — Machine Learning Engineer

Jul 2009 – Mar 2010 | Fostertown (Remote)

- Cut inference cost by 15% through quantisation, batching and autoscaling on Grafana.

Education

High School

Matthewfurt High School | 2009

Technical Skills

Frameworks: NumPy, PyTorch, Scikit-learn, Pandas

Cloud & DevOps: CI/CD, Grafana

Data: Kafka, Databricks, BigQuery, Snowflake, Airflow, Spark, Redshift, MySQL

Testing & Tooling: Cypress

Side Projects

Technical blog

Writing up Kafka experiments and failures, roughly 3k readers a month.

Languages

English – Bilingual

Hindi – C1

Portuguese – Fluent

References available upon request.