

Michael Walker

SUMMARY

27 years in Data/AI/ML, most of it spent on systems other teams depend on. Comfortable leading technical decisions, pushing back on scope when it matters, and carrying a pager for what I build. Primary tools: Databricks and Kafka.

ACADEMIC BACKGROUND

High School

Hannahport High School | 1993 – 1997

RELEVANT EXPERIENCE

Lead Machine Learning Engineer, Ironwood Labs

Jul 2015 – Current | Port Cherylburch (On-site)

- » Productionised demand forecasting models that reduced stock-outs by 22% across 15 regions.
- » Lifted recommendation click-through by 18% through feature engineering and retraining cadence in Spark.
- » Ran A/B tests across 3 model variants; the winning variant added an estimated \$30k in annual revenue.
- » Reduced label noise by redesigning the annotation guidelines, improving validation accuracy by 18%.
- » Mentored 8 analysts moving into modelling work, with weekly paper reading and code review.

Senior Applied Scientist, IBM

Nov 2007 – May 2015 | Christopherburgh

- » Built an end-to-end training pipeline on NumPy, cutting retraining from days to 20 hours.
- » Cut inference cost by 45% through quantisation, batching and autoscaling on Nginx.
- » Established offline/online evaluation parity, catching 15 silent regressions before release.
- » Wrote the team's model documentation and review standard, now required before any production launch.

Machine Learning Engineer, Shopify

Jan 2002 – Jun 2007 | East William (On-site)

- » Translated ambiguous business questions into a modelling roadmap with product and finance stakeholders.
- » Built the feature store that 6 downstream models now read from, ending duplicated feature logic.
- » Migrated ad-hoc notebooks into versioned Pandas pipelines with reproducible runs and tracked experiments.

Summit Digital — MLOps Engineer

Jan 1999 – Aug 2001 | Port Sheilafurt

- » Developed an NLP classification system in TensorFlow reaching 92% F1, replacing a manual review queue.
- » Trained and deployed Pandas models serving 2M predictions per day at under 180ms p95.
- » Led a team of 8 through a full replatforming with no missed delivery dates.

Junior Data Scientist

Harbor Point Tech | Nov 1997 – Dec 1998 | West John (Hybrid - 2 days on-site)

- » Introduced written design docs and an RFC process, cutting rework late in delivery.
- » Ran retrospectives and drove the follow-up actions to completion, not just to the board.
- » Presented the team's work at an internal engineering forum of 200+ colleagues.

SKILLS

- » Databricks
- » Kafka
- » Scikit-learn
- » Pandas
- » Airflow
- » TensorFlow
- » LangChain
- » NumPy
- » Docker
- » SQL

- » Spark
- » Excel
- » Redshift
- » Nginx
- » Node.js
- » Kotlin

NOTABLE PROJECTS

Commuter optimiser

Pulls transit APIs and suggests departure times, built on Nginx.

Conference talk

Spoke on Nginx in production at a regional meetup, ~60 attendees.

CERTIFICATIONS

Certified Information Systems Security Professional (CISSP) — ISC2, 2018

Google Analytics Certification — Google, 2024

RESEARCH OUTPUT

"Scheduling strategies for distributed task queues" - accepted at a peer-reviewed workshop, 1999.

"Scheduling strategies for distributed task queues" - technical report, 1997. Cited internally as the basis for the current design.

PERSONAL INTERESTS

Long-distance swimming, Open-source contribution, Five-a-side football