

Andrew Marks

kimberly35@example.net 686-738-7829 Davidsonville github.com/andrewdev

ABOUT ME

22 years in Data/AI/ML, most of it spent on systems other teams depend on. Comfortable leading technical decisions, pushing back on scope when it matters, and carrying a pager for what I build. Primary tools: Databricks and Kafka.

PROFESSIONAL EXPERIENCE

Vantage Analytics — Senior Data Scientist

02/2014 - Present | North Sarahborough

- Built an end-to-end training pipeline on Python, cutting retraining from days to 4 hours.
- Developed an NLP classification system in NumPy reaching 82% F1, replacing a manual review queue.
- Trained and deployed Python models serving 3M predictions per day at under 80ms p95.
- Wrote the team's model documentation and review standard, now required before any production launch.

Applied Scientist

Foxglove Media | 10/2008 - 10/2013 | Hebertland

- Established offline/online evaluation parity, catching 8 silent regressions before release.
- Built the feature store that 4 downstream models now read from, ending duplicated feature logic.
- Migrated ad-hoc notebooks into versioned MLflow pipelines with reproducible runs and tracked experiments.

Data Scientist

Summit Digital | 05/2006 - 06/2008 | Owenton (Hybrid)

- Productionised demand forecasting models that reduced stock-outs by 25% across 12 regions.
- Translated ambiguous business questions into a modelling roadmap with product and finance stakeholders.

Atlassian — AI Researcher

08/2003 - 01/2006 | Angelachester

- Lifted recommendation click-through by 35% through feature engineering and retraining cadence in MLflow.
- Ran A/B tests across 40 model variants; the winning variant added an estimated \$400k in annual revenue.

EDUCATION

Master, Business Administration

University of Manchester | 2001 - 2003

Magna Cum Laude

Thesis: "Organisational change during platform migration"

Bachelor, Business Administration

Halloway Institute of Technology | 1998 - 2001

CORE COMPETENCIES

Databricks (2 years), Kafka (8 years), Snowflake (12 years), BigQuery (5 years), Python (22 years), Docker (7 years), NumPy (7 years), MLflow (11 years), AWS (2 years), SQL (20 years), TensorFlow (3 years), Bash (4 years), Postman (5 years), Linux (14 years), Nginx (8 years), Grafana (2 years)

OPEN SOURCE

Match-day predictor

Small model in SQL predicting league results. Beats the naive baseline; still loses to the bookmakers.

HONOURS & AWARDS

Internal Hackathon Winner, 2024

LANGUAGES

English – C2