

Michelle Smith

khill@example.com – 483.658.6812x25957 – Sandraland – github.com/michelledev

Summary

9 years in Data/AI/ML, most of it spent on systems other teams depend on. Comfortable leading technical decisions, pushing back on scope when it matters, and carrying a pager for what I build. Primary tools: SQL and PyTorch.

Career History

Senior MLOps Engineer, Cobalt Systems

2023-07 — Present | Lake Danielfort (On-site)

- Built the feature store that 12 downstream models now read from, ending duplicated feature logic.
- Productionised demand forecasting models that reduced stock-outs by 25% across 30 regions.
- Cut inference cost by 12% through quantisation, batching and autoscaling on Grafana.
- Lifted recommendation click-through by 35% through feature engineering and retraining cadence in Pandas.
- Built an end-to-end training pipeline on LangChain, cutting retraining from days to 5 hours.

Vantage Analytics — Machine Learning Engineer

2019-12 — 2023-05 | Kathleenbury

- Established offline/online evaluation parity, catching 30 silent regressions before release.
- Translated ambiguous business questions into a modelling roadmap with product and finance stakeholders.
- Trained and deployed Databricks models serving 40M predictions per day at under 140ms p95.
- Mentored 20 analysts moving into modelling work, with weekly paper reading and code review.

AI Researcher

Nova Tech Group | 2016-09 — 2019-07 | East Jessicafurt

- Developed an NLP classification system in Pandas reaching 92% F1, replacing a manual review queue.
- Ran A/B tests across 15 model variants; the winning variant added an estimated \$150k in annual revenue.
- Migrated ad-hoc notebooks into versioned Databricks pipelines with reproducible runs and tracked experiments.

Education

Bachelor, Economics

Pinehurst University | 2012 — 2016

Graduated with Distinction

Technologies

SQL (1 year), PyTorch (2 years), Snowflake (2 years), Pandas (3 years), MLflow (4 years), LangChain (4 years), Databricks (1 year), Sketch (1 year), Grafana (6 years), Confluence (4 years)

Personal Projects

Commute optimiser

Pulls transit APIs and suggests departure times, built on LangChain.

Match-day predictor

Small model in Pandas predicting league results. Beats the naive baseline; still loses to the bookmakers.