

Charles Bradley

About Me

Pragmatic Data Scientist (3 years in Data/AI/ML). I enjoy untangling messy systems and leaving them simpler than I found them. Day to day that means MLflow, Spark and a lot of conversations with the people who use what we build.

Key Achievements

- » Trained and deployed Spark models serving 2M predictions per day at under 60ms p95.
- » Migrated ad-hoc notebooks into versioned LangChain pipelines with reproducible runs and tracked experiments.
- » Mentored 4 analysts moving into modelling work, with weekly paper reading and code review.

Education and Training

Bachelor, Economics

New Dennisfurt Institute of Technology | 2021

First Class Honours

Employment History

Ironwood Consulting — Data Scientist

May '24 – Present

- » Built an end-to-end training pipeline on MLflow, cutting retraining from days to 6 hours.
- » Lifted recommendation click-through by 15% through feature engineering and retraining cadence in LangChain.
- » Cut inference cost by 22% through quantisation, batching and autoscaling on Docker.

Applied Scientist, Fenwick Digital

Jun '23 – May '24

- » Developed an NLP classification system in LangChain reaching 96% F1, replacing a manual review queue.
- » Ran A/B tests across 50 model variants; the winning variant added an estimated \$500k in annual revenue.

Technical Proficiencies

Languages: Swift, SQL

Frameworks: LangChain

Cloud & DevOps: Docker

Data: MLflow, Spark, Airflow

Testing & Tooling: GraphQL

Notable Projects

Commute optimiser

Pulls transit APIs and suggests departure times, built on SQL.

References available upon request.