

DANIEL BIRD

jessicamiller@example.com 6497386204 West Stevenland linkedin.com/in/daniel-bird

Senior Data/AI/ML professional, 13 years in. Track record of taking ambiguous problems through design, delivery and the unglamorous operational work that follows. I mentor deliberately and write things down. Core stack: BigQuery, Pandas, Python.

RELEVANT EXPERIENCE

Data Scientist

Datadog | Sep '17 – Present

- Lifted recommendation click-through by 35% through feature engineering and retraining cadence in TensorFlow.
- Built the feature store that 50 downstream models now read from, ending duplicated feature logic.
- Established offline/online evaluation parity, catching 4 silent regressions before release.
- Translated ambiguous business questions into a modelling roadmap with product and finance stakeholders.

AI Researcher

Spotify | Mar '14 – Sep '17

- Wrote the team's model documentation and review standard, now required before any production launch.
- Built an end-to-end training pipeline on TensorFlow, cutting retraining from days to 15 hours.
- Trained and deployed NumPy models serving 20M predictions per day at under 80ms p95.

Lumen Works — Junior AI Researcher

Apr '13 – Dec '13

- Ran A/B tests across 40 model variants; the winning variant added an estimated \$400k in annual revenue.
- Mentored 15 analysts moving into modelling work, with weekly paper reading and code review.

Machine Learning Engineer, Cobalt Systems

Dec '12 – Apr '13

- Productionised demand forecasting models that reduced stock-outs by 15% across 30 regions.
- Developed an NLP classification system in Hugging Face reaching 92% F1, replacing a manual review queue.

ACADEMIC HISTORY

Master, Economics

Stanford University | 2009 – 2011

Bachelor, Mathematics

Georgia Institute of Technology | 2009

SKILLS & TOOLS

Languages: Python

Frameworks: Pandas, NumPy, Hugging Face, TensorFlow, React, Django

Cloud & DevOps: AWS, Docker

Data: BigQuery, Snowflake

Testing & Tooling: Confluence, Playwright

PROJECTS

Technical blog

Writing up BigQuery experiments and failures, roughly 8k readers a month.