

MARTHA GILBERT

SUMMARY

Data/AI/ML specialist with 6 years+ building and shipping production systems. Deep hands-on experience with Snowflake, Docker and Airflow, most recently at Redwood Data.

EDUCATION

PhD, Mathematics

Pinehurst University | 2013–2018

Graduated with Distinction

Thesis: “Graph partitioning under capacity constraints”

Bachelor, Physics

Marchmont University | 2009–2012

TECHNOLOGIES

Skill	Experience
Snowflake	3 yrs
Docker	4 yrs
Airflow	2 yrs
Kafka	3 yrs
Git	1 year
Ansible	3 yrs
Bash	1 year

RELEVANT EXPERIENCE

MLOps Engineer, Redwood Data

2024–Present

- Cut inference cost by 70% through quantisation, batching and autoscaling on Docker.
- Productionised demand forecasting models that reduced stock-outs by 55% across 6 regions.
- Developed an NLP classification system in the modelling stack reaching 96% F1, replacing a manual review queue.
- Trained and deployed the modelling stack models serving 15M predictions per day at under 80ms p95.

Keystone Apps — Data Scientist

2022–2024

- Lifted recommendation click-through by 25% through feature engineering and retraining cadence in the modelling stack.
- Established offline/online evaluation parity, catching 12 silent regressions before release.
- Translated ambiguous business questions into a modelling roadmap with product and finance stakeholders.

Machine Learning Engineer, Spotify

2020–2022

- Built an end-to-end training pipeline on the modelling stack, cutting retraining from days to 2 hours.
- Ran A/B tests across 8 model variants; the winning variant added an estimated \$80k in annual revenue.

PROJECTS

Match-day predictor

Small model in Snowflake predicting league results. Beats the naive baseline; still loses to the bookmakers.

Commute optimiser

Pulls transit APIs and suggests departure times, built on Kafka.

RECOGNITION

Employee of the Quarter, 2021

LANGUAGES

English – Bilingual

German – Professional working proficiency

Arabic – Fluent