

Patrick Ferrell

mwilliams@example.org
(400)674-8597
linkedin.com/in/patrick-ferrell

CAREER OBJECTIVE

Senior Data/AI/ML professional, 19 years in. Track record of taking ambiguous problems through design, delivery and the unglamorous operational work that follows. I mentor deliberately and write things down. Core stack: TensorFlow, Docker, Snowflake.

ACADEMIC HISTORY

Master, Information Systems

Oakmere University | 2004–2006

Bachelor, Data Science

Copperfield State University | 2003

KEY SKILLS

TensorFlow (6 yrs), Docker (4 yrs), Snowflake (18 yrs), Databricks (5 yrs), Hugging Face (19 yrs), NumPy (10 yrs), Kafka (6 yrs), AWS (2 yrs), LangChain (4 yrs), SQL (19 yrs), BigQuery (6 yrs), Figma (12 yrs), Confluence (12 yrs), Rust (2 yrs), Angular (13 yrs), FastAPI (12 yrs)

PROFESSIONAL EXPERIENCE

Principal Data Scientist, Uber

2019–Present | New Steven

- » Developed an NLP classification system in TensorFlow reaching 94% F1, replacing a manual review queue.
- » Ran A/B tests across 15 model variants; the winning variant added an estimated \$150k in annual revenue.
- » Trained and deployed LangChain models serving 30M predictions per day at under 110ms p95.

MLOps Engineer

Redwood Data | 2011–2018 | East Chelsealand (Hybrid)

- » Built the feature store that 20 downstream models now read from, ending duplicated feature logic.
- » Built an end-to-end training pipeline on NumPy, cutting retraining from days to 50 hours.

Beacon Interactive — MLOps Engineer

2008–2011 | Matthewton

- » Productionised demand forecasting models that reduced stock-outs by 70% across 5 regions.

Data Scientist, Vantage Analytics

2007–2008 | Ericahaven

- » Cut inference cost by 45% through quantisation, batching and autoscaling on AWS.

AI Researcher, Meridian Software

2006–2006 | Lake Sarahhaven

- » Reduced label noise by redesigning the annotation guidelines, improving validation accuracy by 70%.

OPEN SOURCE

Conference talk

Spoke on TensorFlow in production at a regional meetup, ~30 attendees.

Personal finance tracker

Self-hosted spend tracker built with Hugging Face. Deliberately kept the data local rather than using a hosted service.

References available upon request.